

No. of CPU = 1

output of testcow.c: (After implementing copy-on-write)

Parent start: no. of free page = 56776  
Child1: no. of free page just after fork = 56674  
Child1: no. of free page after some write = 56673  
Child2: no. of free page just after fork = 56708  
Child2: no. of free page after some write = 56707  
Parent end: no. of free page = 56776

output of testcow.c: (Without implementing copy-on-write)

Parent start: no. of free page = 56776  
Child1: no. of free page just after fork = 56636  
Child1: no. of free page after some write = 56636  
Child2: no. of free page just after fork = 56636  
Child2: no. of free page after some write = 56636  
Parent end: no. of free page = 56776

No. of CPU = 2

output of testcow.c:(After implementing copy-on-write)

Parent start: no. of free page = 56775  
Child1: no. of free page just after fork = 56639  
Child1: no. of free page after some write = 56638  
Child2: no. of free page just after fork = 56638  
Child2: no. of free page after some write = 56637  
Parent end: no. of free page = 56775

output of testcow.c:(without implementing copy-on-write)

Parent start: no. of free page = 56775  
Child1: no. of free page just after fork = 56635  
Child1: no. of free page after some write = 56635  
Child2: no. of free page just after fork = 56635  
Child2: no. of free page after some write = 56635  
Parent end: no. of free page = 56775

Explanation:

Without implementing copy-on-write, all the pages needed by new process are allocated at time of fork itself, thus there is no change in number of free page just after the fork statement, and after writing in memory, in case of both child1 and child2.

However, after implementing, copy-on-write, at time of fork, new pages are allocated, only at time of write, at the time of fork no user memory is written, thus there is a difference in number of free page, just after fork and after something is written into user part of memory image of the process, both in the case of child1 and child2.

After both child exits, the number of free page at the start of parent process = no. Of free page at the end of process just before exiting.

## Changes Done:

1. `kalloc.c`
  - a.) defined new variable `numfreepages` to keep track of number of freepage
  - b.) defined an array `pg_ref_array` to keep track of reference count corresponding to each physical location
  - c.) defined a new function `pgcountinit` to be called from `main.c`, to initialize above two variables
  - d.) wrote new `kfree` and `kalloc` function, which modifies `numfreepages` and reference count accordingly
  - e.) defined a struct `pg_ref` for using lock
  - f.) defined function `increase_ref_count(uint location)`, to increase reference count of that location
  - g.) defined function `decrease_ref_count(uint location)`, to decrease reference count of that location
  - h.) defined function `set_ref_count_to_one(uint location)`, to set reference count = 1 of that location
2. `main.c`
  - a.) call function `pginitcount` to initialize at start
3. `user.h`
  - a.) defined function `getNumFreePages()`, which can be called externally
4. `vm.c`
  - a.) wrote new `copyuvm` to implement copy-on-write and `handle_pgflt` function to handle pagefault
5. `testcow.c`
  - a.) test case file
6. `trap.c`
  - a.) trap handler to handle `T_PGFLT` and call `handle_pgflt` function
7. `sysproc.c`
  - a.) defined system call `sys_getNumFreePages`
8. `syscall.h`
  - a.) system call number corresponding to `SYS_getNumFreePages`
9. `syscall.c`
  - a.) defined `SYS_getNumFreePages = sys_getNumFreePages`
10. `defs.h`
  - a.) defined all variable and functions so that they can be accessed from other files as well
11. `usys.S`
  - a.) defined `SYSCALL(getNumFreePages)`
12. `Makefile`
  - a.) added `_testcow`, to compile `testcow.c`
  - b.) `CPUS` changed accordingly to test one 1 or 2 cpus

## Copy-on-Write Implementation

fork() system call calls copyvm() to make memory image of child process, which in turn call setupvm() function to set up kernel parts of memory images, and then originally kalloc was called for getting new page for each user part of memory image and new page is given to the child at the time of fork itself.

However, in the modified copyvm(), only setupkvm() is called. For user part of memory image, new page table entry is made for child process, and the page table entry corresponds to same physical address pointed by the parent, rather than allocating new page using kalloc(). The permission of parent page table entry is set to NOT WRITABLE, and same permission(also NOT WRITABLE), is set for child's parent's table, and the reference count of corresponding physical address is increased by 1, and tlb is flushed and pagetable is reinstalled.

Now, whenever a process wants to write and the permission of corresponding page table entry is NOT WRITABLE, a trap is generated, which in turn call handle\_pgfilt(). handle\_pgfilt() allocates new page using kalloc(), and set its permission to WRITABLE, and decrease the reference count of the old physical address by 1, and the page table is reinstalled.

Since kalloc() always allocates a new page, in kalloc function(), it sets the reference count of that physical address to be 1, and decreases the number of free pages by 1.

kfree() is also called on the physical address which cause the page fault. The modified kfree, first decreases its reference count by 1, and if the reference count goes to 0, it frees up that page as done by kfree() earlier, and increases the number of free pages by 1.

number of free page is initialized with 0, and number of reference count for all physical location by 1, since main function will call kinit1() and kinit2() for free pages, which will eventually call kfree(), and the reference count will be set to 0 there.